



**University of Tehran**  
**Electrical and Computer Engineering Department**  
**Neural Networks and Deep Learning**  
**Extra Homework**

Name	Danial Saeedi
Student No	810198571
Date	Wednesday, December 15, 2021

## Table of contents

1. Solving CartPole Problem using Policy Gradient Method .....	3
--	---

# 1. Solving CartPole Problem using Policy Gradient Method

## Part 1

### Import Dependencies

First of all, we have to import dependencies:

```
In [ ]: import numpy as np
import torch.nn as nn
import torch.nn.functional as F
import torch
import gym
from torch.autograd import Variable
import random
```

### Define Parameters

Through trial and error, the chosen learning rate and gamma parameters are 0.01 and 0.99 respectively.

```
In [ ]: LR = 0.01
GAMMA = 0.99
```

### Creating an environment in OpenAI Gym

```
In [ ]: env = gym.make('CartPole-v0').unwrapped
history = []
```

## Setting up the policy network

The input and output size of policy network is 4 and 2 respectively because the dimension of observation space is 4 and there are 2 possible actions(either going left or right) in this problem. The chosen optimizer is Adam.

```
In [ ]: class Network(nn.Module):
        def __init__(self):
            super(Network, self).__init__()
            # define forward pass with one hidden layer with ReLU activation and softmax after output layer
            self.l1 = nn.Linear(4, 150)
            self.l2 = nn.Linear(150, 2)
        def forward(self, x):
            x = F.relu(self.l1(x))
            x = F.softmax(self.l2(x))
            return x

In [ ]: model = Network()

use_cuda = torch.cuda.is_available()
if use_cuda:
    model.cuda()
FloatTensor = torch.cuda.FloatTensor if use_cuda else torch.FloatTensor
LongTensor = torch.cuda.LongTensor if use_cuda else torch.LongTensor

optim = torch.optim.Adam(model.parameters(), lr=LR)
```

## Calculate Discount Rewards

Here, we calculate discount reward based on this formula:

$$R_t = \sum_{k=t}^T \gamma^{(k-t)} r_k(s_k, a_k)$$

```
In [ ]: def calculate_discount_rewards(r):
        discounted_r = torch.zeros(r.size())
        running_add = 0
        for t in reversed(range(len(r))):
            running_add = running_add * GAMMA + r[t]
            discounted_r[t] = running_add
        return discounted_r
```

## Training

The summary of section:

- Calculate the probability of the action taken at each time step.
- Multiply the probability by the discounted return (the sum of rewards).
- Use this probability-weighted return to backpropagate and minimize the loss.

```
In [ ]: for e in range(10000):
        complete = run_episode(model, e, env)

        if complete:
            break

[Episode 0] reward: 12.0
[Episode 1] reward: 14.0
[Episode 2] reward: 31.0
[Episode 3] reward: 42.0
[Episode 4] reward: 33.0
[Episode 5] reward: 25.0
[Episode 6] reward: 27.0
[Episode 7] reward: 25.0
[Episode 8] reward: 33.0
[Episode 9] reward: 36.0
[Episode 10] reward: 75.0
[Episode 11] reward: 50.0
[Episode 12] reward: 39.0
[Episode 13] reward: 73.0
[Episode 14] reward: 31.0
[Episode 15] reward: 28.0
[Episode 16] reward: 53.0
[Episode 17] reward: 83.0
[Episode 18] reward: 100.0
[Episode 19] reward: 41.0
[Episode 20] reward: 103.0
[Episode 21] reward: 169.0
[Episode 22] reward: 159.0
```

## Run Episode

```
In [ ]: def run_episode(net, e, env):
    state = env.reset()
    reward_sum = 0
    xs = FloatTensor([])
    ys = FloatTensor([])
    rewards = FloatTensor([])
    steps = 0

    while True:
        x = FloatTensor([state])
        xs = torch.cat([xs, x])

        action_prob = net(Variable(x))

        # select an action depends on probability
        action = 0 if random.random() < action_prob.data[0][0] else 1

        y = FloatTensor([[1, 0]] if action == 0 else [[0, 1]])
        ys = torch.cat([ys, y])

        state, reward, done, _ = env.step(action)
        rewards = torch.cat([rewards, FloatTensor([[reward]])])
        reward_sum += reward
        steps += 1

    if done or steps >= 500:
        adv = calculate_discount_rewards(rewards)
        adv = (adv - adv.mean()) / (adv.std() + 1e-7)
        loss = learn(xs, ys, adv)
        history.append(reward_sum)
        print("[Episode {:>5}] reward: {}".format(e, reward_sum))
        if sum(history[-5:])/5 > 490:
            return True
        else:
            return False
```

## Learning Function

```
In [ ]: def learn(x, y, adv):  
    # calculate probabilities of taking each action  
    action_pred = model(Variable(x))  
    y = Variable(y, requires_grad=True)  
    adv = Variable(adv).cuda()  
    log_lik = -y * torch.log(action_pred)  
    log_lik_adv = log_lik * adv  
    loss = torch.sum(log_lik_adv, 1).mean()  
  
    optim.zero_grad()  
    loss.backward()  
    optim.step()  
  
    return loss.data
```

## Plotting average reward per episode

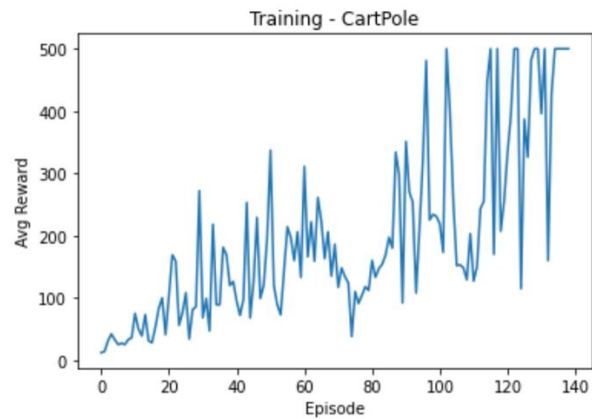


Figure 1 Average reward per episode

## Part 2

These type of problem in RL require techniques such as mathematical programming and heuristic reduction to keep them manageable. A linear program with tons of actions might be solvable within seconds. Lots of real-life problems are convex or even (approximately) linear. This is a very powerful property that often makes problem solving considerably easier.

But linear programming can't solve every problem. The solution proposed here is based on this [paper](#). The paper proposed a new policy architecture called Wolpertinger. This architecture avoids the heavy cost of evaluating all actions while **retaining generalization over actions**. The solution is based on actor-critic framework. We use multi-layer neural networks as function approximators for both our actor and critic functions. Training this policy is based on [Deep Deterministic Policy Gradient](#).

---

**Algorithm 1** Wolpertinger Policy

---

State  $\mathbf{s}$  previously received from environment.  
 $\hat{\mathbf{a}} = f_{\theta^\pi}(\mathbf{s})$  {Receive proto-action from actor.}  
 $\mathcal{A}_k = g_k(\hat{\mathbf{a}})$  {Retrieve  $k$  approximately closest actions.}  
 $\mathbf{a} = \arg \max_{\mathbf{a}_j \in \mathcal{A}_k} Q_{\theta^Q}(\mathbf{s}, \mathbf{a}_j)$   
Apply  $\mathbf{a}$  to environment; receive  $r, \mathbf{s}'$ .

---

Figure 2 Wolpertinger Policy Algorithm

### 1) Action Generation

This architecture reasons over actions within a **continuous space**  $\mathbb{R}^n$ , and then maps this output to the discrete action Discrete Action Spaces set  $\mathcal{A}$ :

$$f_{\theta^\pi} : \mathcal{S} \rightarrow \mathbb{R}^n$$
$$f_{\theta^\pi}(\mathbf{s}) = \hat{\mathbf{a}}.$$

$f_{\theta^\pi}$  is a function parametrized by  $\theta^\pi$ , mapping from the state representation space  $\mathbb{R}^m$  to the action representation space  $\mathbb{R}^n$ . We need to be able to map from  $\hat{\mathbf{a}}$  to an element in  $\mathcal{A}$

$$g : \mathbb{R}^n \rightarrow \mathcal{A}$$
$$g_k(\hat{\mathbf{a}}) = \arg \min_{\mathbf{a} \in \mathcal{A}}^k \|\mathbf{a} - \hat{\mathbf{a}}\|_2.$$

$g_k$  is a  $k$ -nearest-neighbor mapping from a continuous space to a discrete set. It returns the  $k$  actions in  $\mathcal{A}$  that are closest to  $\hat{\mathbf{a}}$  by  $L_2$  distance. This is called proto-action. In the bottom half of Figure 2 we can see proto action.

## 2) Action Refinement

Certain actions may be near each other in the action embedding space, but in certain states they must be distinguished as one has a particularly low long-term value relative to its neighbors.

The model refines the choice of action by selecting the highest-scoring action according to Q-value:

$$\pi_{\theta}(\mathbf{s}) = \arg \max_{a \in g_k \circ f_{\theta} \pi(\mathbf{s})} Q_{\theta Q}(\mathbf{s}, \mathbf{a})$$

The size of the generated action set,  $k$ , is task specific, and allows for an explicit trade-off between policy quality and speed.




Figure 3 Wolpertinger Architecture